

Claxon competition 2024-Copy1

January 4, 2025

```
[1]: <h2>Financial institutions face significant risks due to loan defaults.␣
    ↪Accurately predicting the
probability of default (PD) on loans is critical for risk management and␣
    ↪strategic planning.
In this competition, participants are tasked with developing a predictive model␣
    ↪that estimates the
probability of default on loans using historical loan data.</h2>

### Objective:
The objective is to build a predictive model on this data to help the bank␣
    ↪decide on whether to approve a loan to a prospective applicant.

###Data Dictionary

    unnamed-observation number.
    loan_id - unique identifier for each loan.
    Sex (Categories: male, female, other)
    disbursemet_date-date when loan funds were released and made available to␣
    ↪the borrower.
    currency-currency in which the loan was issued(USD).
    country- country of origin for the borrower.
    sex - gender of the customer.
    is_employed- Not employed(False), employed(True).
    job- the job of the customer.
    location- place of residency of the client.
    loan_amount - amount for which loan is requested.
    number_of_defaults- the count of defaulted times of a customer.
    outstanding_balance- the amount not yet paid by the customer.
    interest_rate-percentage of loan amount that a customer pays to the bank as␣
    ↪interest over a year.
    age - age of the customer.
    remaining term- amount of time remaining for the loan to be fully repaid.
    salary-income of the customer.
    marital_status- 1-married, 0-single.
    Loan Status- 1-defaulted, 0-Did not default.
```

```
Cell In[1], line 1
    <h2>Financial institutions face significant risks due to loan defaults.
↪ Accurately predicting the
    ^
SyntaxError: invalid syntax
```

```
[3]: ### Import necessary libraries
```

```
[5]: # this will help in making the Python code more structured automatically (good
    ↪ coding practice)
import jupyter_black

jupyter_black.load()

# To filter the warnings
import warnings

warnings.filterwarnings("ignore")

# Libraries to help with reading and manipulating data
import pandas as pd
import numpy as np

# Removes the limit for the number of displayed columns
pd.set_option("display.max_columns", None)
# Sets the limit for the number of displayed rows
pd.set_option("display.max_rows", 200)

# libraries to help with data visualization
import matplotlib.pyplot as plt
import seaborn as sns

# Library to split data
from sklearn.model_selection import train_test_split
from sklearn.preprocessing import StandardScaler

# To build model for prediction
import statsmodels.stats.api as sms
from statsmodels.stats.outliers_influence import variance_inflation_factor
import statsmodels.api as sm
from statsmodels.tools.tools import add_constant
from sklearn.linear_model import LogisticRegression # Logistic Regression
from sklearn.tree import DecisionTreeClassifier # Decision Tree
from sklearn.ensemble import (
```

```

    RandomForestClassifier,
    GradientBoostingClassifier,
) # Random Forest, Gradient Boosting Machines
from sklearn.svm import SVC # Support Vector Machines

import pickle

# To get diferent metric scores
from sklearn.metrics import (
    f1_score,
    accuracy_score,
    recall_score,
    precision_score,
    confusion_matrix,
    roc_auc_score,
    ConfusionMatrixDisplay,
    precision_recall_curve,
    roc_curve,
)
import pandas as pd

pd.set_option("display.max_columns", None)

```

<IPython.core.display.HTML object>

```
[7]: # Loading the dataset - sheet_name parameter is used if there are multiple tabs
      ↪ in the excel file.
```

```
[9]: df = pd.read_csv(
      "C:/Users/Munashe Muchinako/OneDrive/Desktop/data science/Data Science_
      ↪Competition Question and Data/data_science_competition_2024.csv"
    )
```

```
[11]: # copy the data into duplicate variable 'data' to avoid making changes to the
      ↪original data
```

```
[13]: data_raw = df.copy()
```

```
[15]: # show top 5 rows in the data
```

```
[17]: data_raw.head(4)
```

```
[17]:
```

	Unnamed: 0	loan_id	gender	disbursemet_date	\
0	0	8d05de78-ff32-46b1-aeb5-b3190f9c158a	female	2022 10 29	
1	1	368bf756-fcf2-4822-9612-f445d90b485b	other	2020 06 06	
2	2	6e3be39e-49b5-45b5-aab6-c6556de53c6f	other	2023 09 29	
3	3	191c62f8-2211-49fe-ba91-43556b307871	female	2022 06 22	

	currency	country	sex	is_employed	job	location	loan_amount \
0	USD	Zimbabwe	female	True	Teacher	Beitbridge	39000.0
1	USD	Zimbabwe	other	True	Teacher	Harare	27000.0
2	USD	Zimbabwe	other	True	Nurse	Gweru	35000.0
3	USD	Zimbabwe	female	True	Doctor	Rusape	24000.0

	number_of_defaults	outstanding_balance	interest_rate	age \
0	0	48653.011473	0.22	37
1	2	28752.062237	0.20	43
2	1	44797.554126	0.22	43
3	0	35681.496413	0.23	47

	number_of_defaults.1	remaining term	salary	marital_status	age.1 \
0	0	47	3230.038869	married	37
1	2	62	3194.139103	single	43
2	1	57	3330.826656	married	43
3	0	42	2246.797020	divorced	47

Loan Status

0	Did not default
1	Did not default
2	Did not default
3	Did not default

```
[19]: # Display last 3 rows of the data
```

```
[21]: data_raw.tail(3)
```

```
[21]: Unnamed: 0          loan_id  gender \
99997      99997  4f10e845-8f75-4cd5-9f3a-3dad3e04a483  female
99998      99998  eded01ca-79d2-4e86-a1e3-2ea1354edca7   male
99999      99999  a37561ec-0901-4350-8a13-634f80ece55d   other
```

	disbursement_date	currency	country	sex	is_employed	job \
99997	2021 10 20	USD	Zimbabwe	female	True	Data Analyst
99998	2021 08 22	USD	Zimbabwe	male	True	Engineer
99999	2022 04 29	USD	Zimbabwe	other	True	Engineer

	location	loan_amount	number_of_defaults	outstanding_balance \
99997	Kadoma	48000.0	0	34266.224130
99998	Mutare	36000.0	2	71546.024917
99999	Gweru	46000.0	0	43141.102930

	interest_rate	age	number_of_defaults.1	remaining term	salary \
99997	0.23	43	0	53	3535.599759
99998	0.22	49	2	59	3082.407123
99999	0.21	47	0	47	2670.766532

	marital_status	age.1	Loan Status
99997	married	43	Did not default
99998	single	49	Did not default
99999	married	47	Did not default

```
[23]: # Understand the data shape
```

```
[25]: data_raw.shape
```

```
[25]: (100000, 21)
```

```
[27]: # There are 100000 observations and 21 columns in the dataset
```

```
[29]: ### Check the data types of the columns in the dataset.
data_raw.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 100000 entries, 0 to 99999
Data columns (total 21 columns):
#   Column                                Non-Null Count  Dtype
---  -
0   Unnamed: 0                            100000 non-null  int64
1   loan_id                               100000 non-null  object
2   gender                                100000 non-null  object
3   disbursement_date                     100000 non-null  object
4   currency                              100000 non-null  object
5   country                               99900 non-null   object
6   sex                                   100000 non-null  object
7   is_employed                           100000 non-null  bool
8   job                                   95864 non-null   object
9   location                              99405 non-null   object
10  loan_amount                           100000 non-null  float64
11  number_of_defaults                     100000 non-null  int64
12  outstanding_balance                    100000 non-null  float64
13  interest_rate                          100000 non-null  float64
14  age                                    100000 non-null  int64
15  number_of_defaults.1                   100000 non-null  int64
16  remaining term                         100000 non-null  object
17  salary                                100000 non-null  float64
18  marital_status                         100000 non-null  object
19  age.1                                 100000 non-null  int64
20  Loan Status                           100000 non-null  object
dtypes: bool(1), float64(4), int64(5), object(11)
memory usage: 15.4+ MB
```

```
[31]: -We have 8 continuous variables(age.1, age, salary, Unnamed, loan_amount,
↳number_of_defaults, outstanding_balance, number_of_defaults.1
```

-All other are categorical
 -We can see that there are missing records in the dataset

Cell In[31], line 1

```
-We have 8 continuous variables(age.1, age, salary, Unnamed, loan_amount,
↳number_of_defaults, outstanding_balance, number_of_defaults.1
```

SyntaxError: invalid syntax

[33]: *###Summary of the data*

[35]: data_raw.describe().T

```
[35]:
```

	count	mean	std	min	\
Unnamed: 0	100000.0	49999.500000	28867.657797	0.0	
loan_amount	100000.0	31120.000000	15895.093631	1000.0	
number_of_defaults	100000.0	0.441970	0.688286	0.0	
outstanding_balance	100000.0	36964.909763	10014.758477	0.0	
interest_rate	100000.0	0.210435	0.018725	0.1	
age	100000.0	43.570690	4.863760	21.0	
number_of_defaults.1	100000.0	0.441970	0.688286	0.0	
salary	100000.0	2781.804324	696.450055	250.0	
age.1	100000.0	43.570690	4.863760	21.0	

	25%	50%	75%	max
Unnamed: 0	24999.750000	49999.500000	74999.250000	99999.0
loan_amount	21000.000000	31000.000000	40000.000000	273000.0
number_of_defaults	0.000000	0.000000	1.000000	2.0
outstanding_balance	29625.227472	35063.852394	42133.388817	150960.0
interest_rate	0.200000	0.210000	0.220000	0.3
age	40.000000	44.000000	47.000000	65.0
number_of_defaults.1	0.000000	0.000000	1.000000	2.0
salary	2273.929349	2665.441567	3146.577655	10000.0
age.1	40.000000	44.000000	47.000000	65.0

[37]: Observations

Mean value for the age column is approx 44 and the median is 44. This shows
 ↳that majority of the customers are under 44 years of age.

Mean loan_amount is approx 31120 but it has a wide range with values from
 ↳1000 to 273000. We will explore this further in univariate analysis.

Mean salary is 2782 and median is approx 2665. This shows that majority of
 ↳the customers earn salaries below 2782.

Mean value for outstanding balance is 36965.

Cell In[37], line 3

Mean value for the age column is approx 44 and the median is 44. This shows,
↳ that majority of the customers are under 44 years of age.

IndentationError: unexpected indent

```
[39]: ###Display number of missing values per each column
```

```
[41]: data_raw.isna().sum()
```

```
[41]: Unnamed: 0          0
      loan_id         0
      gender         0
      disbursemet_date 0
      currency        0
      country        100
      sex            0
      is_employed     0
      job            4136
      location        595
      loan_amount     0
      number_of_defaults 0
      outstanding_balance 0
      interest_rate   0
      age            0
      number_of_defaults.1 0
      remaining term  0
      salary         0
      marital_status  0
      age.1          0
      Loan Status    0
      dtype: int64
```

```
[43]: -The country, job and location variables has some missing values, we will
      ↳ impute them using mode(most frequent value) since there relatively few
      ↳ missing values compared to total dataset size
```

Cell In[43], line 1

-The country, job and location variables has some missing values, we will
↳ impute them using mode(most frequent value) since there relatively few missing
↳ values compared to total dataset size

SyntaxError: invalid syntax

```
[45]: ###Data Cleaning
```

```
[47]: # Handling missing data
```

```
[49]: mode_country = data_raw["country"].mode()[0] # calculate mode for country
mode_job = data_raw["job"].mode()[0] # calculate mode for job
mode_location = data_raw["location"].mode()[0] # calculate mode for location

# Fill missing values with mode
data_raw["country"].fillna(mode_country, inplace=True)
data_raw["job"].fillna(mode_job, inplace=True)
data_raw["location"].fillna(mode_location, inplace=True)
```

```
[51]: # Checking missing data again
data_raw.isnull().values.any()
```

```
[51]: False
```

```
[53]: # We can see that there are no longer missing records
```

```
[55]: # Changing the target column loan status from string datatype to
      ↪ boolean-(defaulted=1 otherwise 0)
data_raw["Loan Status"] = np.where(data_raw["Loan Status"] == "Defaulted", 1, 0)
```

```
[57]: # Check for data row duplication between "age" and "age.1"
duplicate_rows = data_raw["age"].equals(data_raw["age.1"])
# Print the result (True if duplication, False otherwise)
print(duplicate_rows)
```

True

```
[59]: # Check for data row duplication between "gender" and "sex"
duplicate_rows = data_raw["gender"].equals(data_raw["sex"])
# Print the result (True if duplication, False otherwise)
print(duplicate_rows)
```

True

```
[61]: # The columns age and age.1 are duplicates as shown by the result above
```

```
[63]: # Check for data row duplication between "number_of_defaults" and
      ↪ "number_of_defaults.1"
duplicate_rows = data_raw["number_of_defaults"].
      ↪ equals(data_raw["number_of_defaults.1"])
# Print the result (True if duplication, False otherwise)
print(duplicate_rows)
```

True


```
[65]: # The columns number_of_defaults and number_of_defaults.1 are duplicates as
      ↪shown by the result above
```

```
[67]: # Remove the duplicated columns, leaving original column and remove the unnamed
      ↪column replace 'Unnamed' with the actual label)
data_raw.drop(
    ["age.1", "number_of_defaults.1", "sex", data_raw.columns[0]], axis=1,
    ↪inplace=True
)
```

```
[69]: # Convert 'date' column to datetime from string
data_raw["disbursemnet_date"] = pd.to_datetime(data_raw["disbursemnet_date"])
```

```
[71]: ##converting remaining term from string type to float
data_raw["remaining term"] = data_raw["remaining term"].str.replace("_", "")
data_raw["remaining term"] = data_raw["remaining term"].astype(float)
```

```
[73]: # check for duplicate rows based on all columns
duplicate_rows = data_raw[data_raw.duplicated()]
print(duplicate_rows)
```

Empty DataFrame

Columns: [loan_id, gender, disbursemnet_date, currency, country, is_employed, job, location, loan_amount, number_of_defaults, outstanding_balance, interest_rate, age, remaining term, salary, marital_status, Loan Status]
Index: []

```
[75]: # There are no more duplicates
```

```
[77]: # Taking a closer look on currency column
data_raw["currency"].value_counts()
```

```
[77]: currency
USD      99980
$USD      20
Name: count, dtype: int64
```

```
[79]: # Formatting value $USD to USD
data_raw["currency"] = data_raw["currency"].str.replace("$USD", "USD")
data_raw["currency"].value_counts()
```

```
[79]: currency
USD      100000
Name: count, dtype: int64
```

```
[81]: # getting rid of empty spaces before each location
data_raw["location"] = data_raw["location"].str.strip()
data_raw["location"].value_counts()
```

```
[81]: location
      Harare          9148
      Bulawayo       8263
      Mutare         8262
      Gweru          7983
      Masvingo       7665
      Marondera      7513
      Rusape         6506
      Chivhu         6411
      Plumtree       5552
      Beitbridge     5311
      Chipinge       4447
      Chimanimani    4388
      Kwekwe         3491
      Chiredzi       3199
      Kadoma         3118
      Nyanga         2142
      Karoi          1899
      Shurugwi       1359
      Zvishavane     1301
      Gokwe          920
      Kariba         671
      Victoria Falls 219
      Redcliff       191
      Hwange         41
      Name: count, dtype: int64
```

```
[83]: # checking counts of data values in the country column
      data_raw["country"].value_counts()
```

```
[83]: country
      Zimbabwe      99887
      zimbabwe       100
      Zim            13
      Name: count, dtype: int64
```

```
[85]: # Correcting the all the values to become Zimbabwe
      data_raw["country"] = data_raw["country"].str.title()
      data_raw["country"] = data_raw["country"].str.replace("Zim", "Zimbabwe")
      data_raw["country"] = data_raw["country"].str.replace("Zimbabwebabwe", "Zimbabwe")
      data_raw["country"].value_counts()
```

```
[85]: country
      Zimbabwe      100000
      Name: count, dtype: int64
```

```
[87]: # Taking a closer look on various customer jobs
data_raw["job"].value_counts()
```

```
[87]: job
Engineer          20660
Nurse             15284
Data Analyst      13204
Doctor            12186
Software Developer 11932
Teacher           8950
Accountant         7802
SoftwareDeveloper  3564
Data Scientist    3521
Lawyer             2862
Data Scintist      35
Name: count, dtype: int64
```

```
[89]: ##Correcting typing errors in job values
data_raw["job"] = data_raw["job"].replace("Data Scintist", "Data Scientist")
data_raw["job"] = data_raw["job"].replace("SoftwareDeveloper", "Software_
Developer")
```

```
[91]: data_raw.head()
```

```
[91]:
```

	loan_id	gender	disbursemet_date	currency	\
0	8d05de78-ff32-46b1-aeb5-b3190f9c158a	female	2022-10-29	USD	
1	368bf756-fcf2-4822-9612-f445d90b485b	other	2020-06-06	USD	
2	6e3be39e-49b5-45b5-aab6-c6556de53c6f	other	2023-09-29	USD	
3	191c62f8-2211-49fe-ba91-43556b307871	female	2022-06-22	USD	
4	477cd8a1-3b01-4623-9318-8cd6122a8346	male	2023-02-08	USD	

	country	is_employed	job	location	loan_amount	\
0	Zimbabwe	True	Teacher	Beitbridge	39000.0	
1	Zimbabwe	True	Teacher	Harare	27000.0	
2	Zimbabwe	True	Nurse	Gweru	35000.0	
3	Zimbabwe	True	Doctor	Rusape	24000.0	
4	Zimbabwe	True	Nurse	Chipinge	19000.0	

	number_of_defaults	outstanding_balance	interest_rate	age	\
0	0	48653.011473	0.22	37	
1	2	28752.062237	0.20	43	
2	1	44797.554126	0.22	43	
3	0	35681.496413	0.23	47	
4	0	34156.055882	0.20	42	

	remaining term	salary	marital_status	Loan Status
0	47.0	3230.038869	married	0

1	62.0	3194.139103	single	0
2	57.0	3330.826656	married	0
3	42.0	2246.797020	divorced	0
4	45.0	2310.858441	married	0

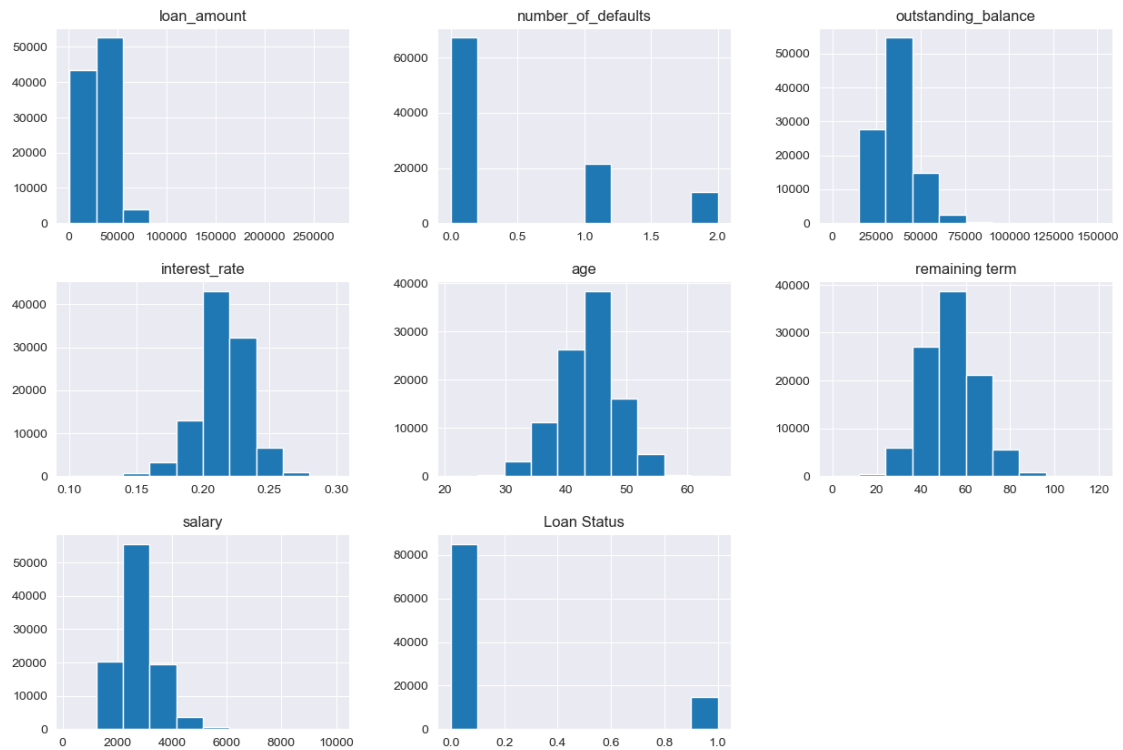
```
[93]: data_raw.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 100000 entries, 0 to 99999
Data columns (total 17 columns):
#   Column                Non-Null Count  Dtype
---  -
0   loan_id                100000 non-null object
1   gender                 100000 non-null object
2   disbursement_date      100000 non-null datetime64[ns]
3   currency               100000 non-null object
4   country                100000 non-null object
5   is_employed            100000 non-null bool
6   job                   100000 non-null object
7   location               100000 non-null object
8   loan_amount            100000 non-null float64
9   number_of_defaults     100000 non-null int64
10  outstanding_balance     100000 non-null float64
11  interest_rate           100000 non-null float64
12  age                     100000 non-null int64
13  remaining term          100000 non-null float64
14  salary                  100000 non-null float64
15  marital_status          100000 non-null object
16  Loan Status             100000 non-null int32
dtypes: bool(1), datetime64[ns](1), float64(5), int32(1), int64(2), object(7)
memory usage: 11.9+ MB
```


```
[95]: data = data_raw.copy()
```

```
[97]: ##### Distribution of variables in the data
```

```
[99]: ## Univariate analysis
sns.set_style("darkgrid")
data.iloc[:, 8:17].hist(figsize=(15, 10))
plt.show()
```



```
[100]: data.iloc[:, 8:17].boxplot(figsize=(20, 5))
plt.show()
```



```
[102]: # Checking the loan Status distribution among defaults and non defaults
data["Loan Status"].value_counts(1)
```

```
[102]: Loan Status
0      0.85134
1      0.14866
Name: proportion, dtype: float64
```

```
[105]: # It can be clearly seen that there is only 14.9% of defaulters, compared to
      ↪ 85% who did not default
```

```
[107]: # Cross-table of gender and Loan Status
round(data.groupby(["gender"])["Loan Status"].value_counts(1), 2).unstack()
```

```
[107]: Loan Status      0      1
gender
female      0.88  0.12
male        0.84  0.16
other       0.84  0.16
```

```
[109]: # It seems that among gender the proportion of women who defaults is lower,
      ↪ compared to the other gender
```

```
[111]: # check if any value in each column is zero
has_zeros_in_loan_amount = (data["loan_amount"] == 0).any()
print(has_zeros_in_loan_amount)
```

False

```
[113]: # There is no observations with zero values
```

```
[115]: def histogram_boxplot(data, feature, figsize=(12, 7), kde=False, bins=None):
      """
      Boxplot and histogram combined

      data: dataframe
      feature: dataframe column
      figsize: size of figure (default (12,7))
      kde: whether to show the density curve (default False)
      bins: number of bins for histogram (default None)
      """

      f2, (ax_box2, ax_hist2) = plt.subplots(
          nrows=2, # Number of rows of the subplot grid= 2
          sharex=True, # x-axis will be shared among all subplots
          gridspec_kw={"height_ratios": (0.25, 0.75)},
          figsize=figsize,
      ) # creating the 2 subplots
      sns.boxplot(
          data=data, x=feature, ax=ax_box2, showmeans=True, color="violet"
      ) # boxplot will be created and a star will indicate the mean value of the
      ↪ column
      sns.histplot(
          data=data, x=feature, kde=kde, ax=ax_hist2, bins=bins, palette="winter"
      ) if bins else sns.histplot(
          data=data, x=feature, kde=kde, ax=ax_hist2
```

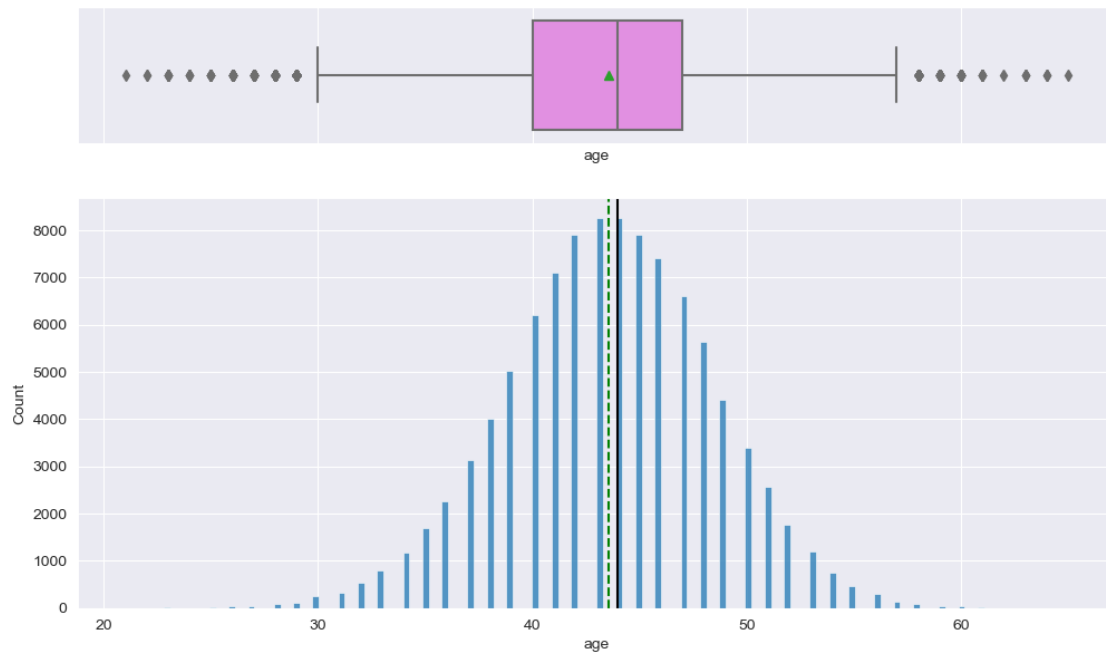
```

) # For histogram
ax_hist2.axvline(
    data[feature].mean(), color="green", linestyle="--"
) # Add mean to the histogram
ax_hist2.axvline(
    data[feature].median(), color="black", linestyle="--"
) # Add median to the histogram

```

```
[117]: ### Observation on Age
```

```
[119]: histogram_boxplot(data, "age")
```



```
[120]: - The distribution of age is equal
- The boxplot shows that there are outliers at both ends
- We will not treat these outliers as they represent the real market trend
```

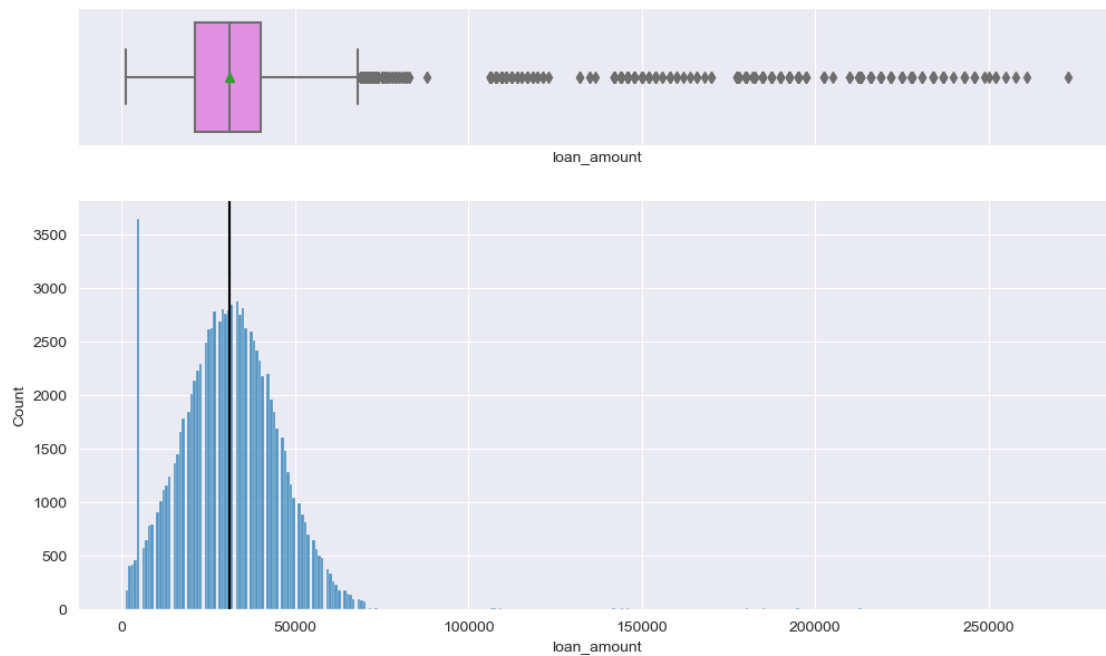
Cell In[120], line 1

```
- The distribution of age is equal
```

SyntaxError: invalid syntax

```
[122]: ### Observation on Credit Amount
```

```
[124]: histogram_boxplot(data, "loan_amount")
```



```
[125]: The distribution of the loan_amount is right-skewed
The boxplot shows that there are outliers at the right end
We will not treat these outliers as they represent the real market trend
```

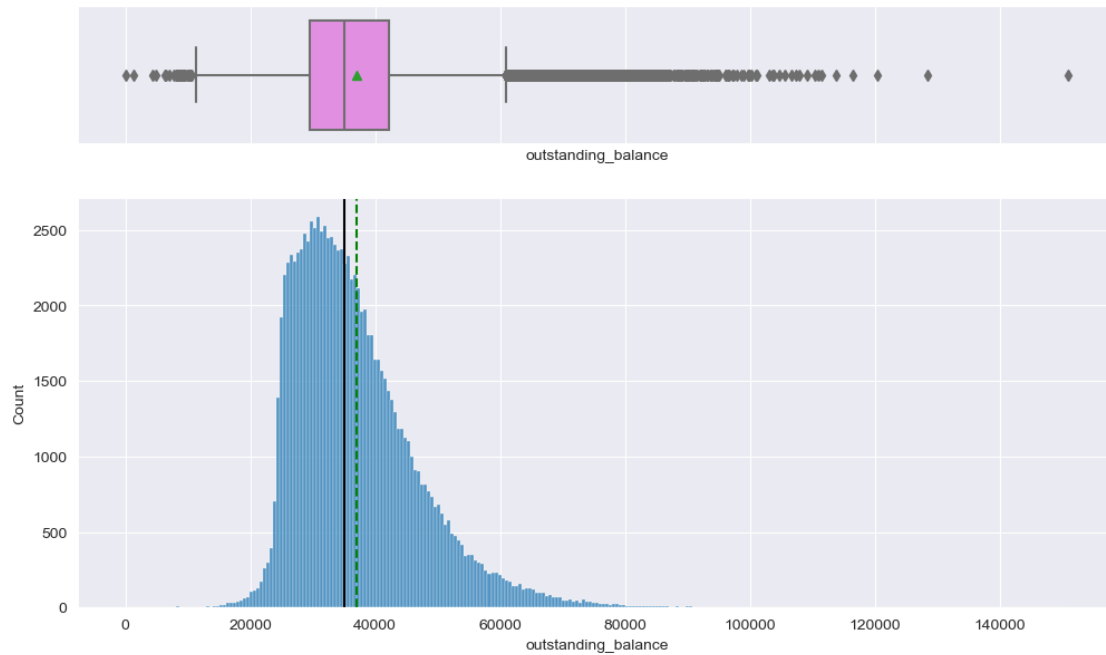
```
Cell In[125], line 1
```

```
The distribution of the loan_amount is right-skewed
```

```
SyntaxError: invalid syntax
```

```
[ ]: ### Observations on Duration
```

```
[128]: histogram_boxplot(data, "outstanding_balance")
```

[129]: The distribution of the outstanding_balance is right-skewed
 The boxplot shows that there are outliers at both ends
 We will not treat these outliers as they represent the real market trend

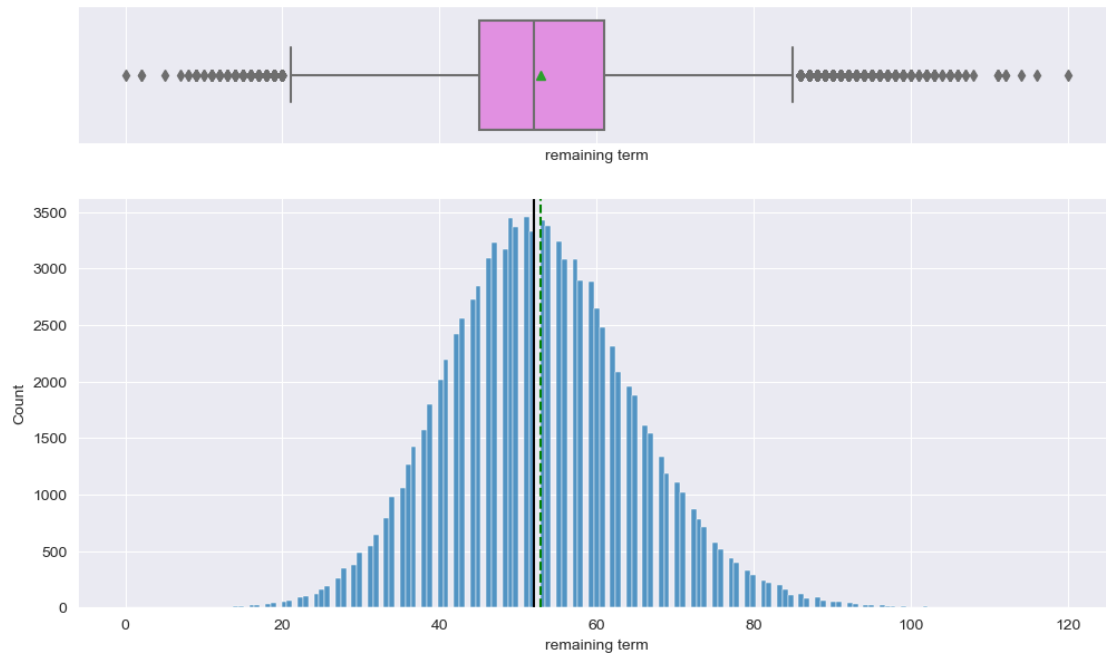
Cell In[129], line 1

The distribution of the outstanding_balance is right-skewed

SyntaxError: invalid syntax

[131]: *### Observations on remaining term for the loan to be fully repaid*

[133]: histogram_boxplot(data, "remaining term")



[135]: -The remaining time is equally distributed

Cell In[135], line 1

-The remaining time is equally distributed

SyntaxError: invalid syntax

[137]: # function to create labeled barplots

```
def labeled_barplot(data, feature, perc=False, n=None):
    """
    Barplot with percentage at the top

    data: dataframe
    feature: dataframe column
    perc: whether to display percentages instead of count (default is False)
    n: displays the top n category levels (default is None, i.e., display all
    ↪ levels)
    """

    total = len(data[feature]) # length of the column
    count = data[feature].nunique()
```

```

if n is None:
    plt.figure(figsize=(count + 1, 5))
else:
    plt.figure(figsize=(n + 1, 5))

plt.xticks(rotation=90, fontsize=15)
ax = sns.countplot(
    data=data,
    x=feature,
    palette="Paired",
    order=data[feature].value_counts().index[:n].sort_values(),
)

for p in ax.patches:
    if perc == True:
        label = "{:.1f}%".format(
            100 * p.get_height() / total
        ) # percentage of each class of the category
    else:
        label = p.get_height() # count of each level of the category

    x = p.get_x() + p.get_width() / 2 # width of the plot
    y = p.get_height() # height of the plot


    ax.annotate(
        label,
        (x, y),
        ha="center",
        va="center",
        size=12,
        xytext=(0, 5),
        textcoords="offset points",
    ) # annotate the percentage

plt.show() # show the plot

```

```
[139]: ### Observations on loan status
```

```
[141]: labeled_barplot(data, "Loan Status", perc=True)
```



```
[143]: - The class distribution in the target variable is imbalanced.
- We have 85.1% observations for non-defaulters and 14.9% observations for
↳ defaulters.
```

```
Cell In[143], line 1
```


```
- The class distribution in the target variable is imbalanced.
```

```
SyntaxError: invalid syntax
```

```
[ ]:
```

```
[146]: ###Observation on sex
```

```
[148]: labeled_barplot(data, "gender", perc=True)
```



```
[150]: - Male customers are taking more credit than female customers
- There are 35% male customers and 32.7% female customers
-The other portion belongs to 'other' class
```


```
Cell In[150], line 1
```

```
- Male customers are taking more credit than female customers
```

```
SyntaxError: invalid syntax
```

```
[152]: ###Observation on marriage status
```

```
[154]: labeled_barplot(data, "marital_status", perc=True)
```



```
[156]: """Majority of the customers i.e. 44% fall into the married category which
↳makes sense as these may be the persons who require loans to help them
↳supply family needs.
There are only approx 27% customers that lie in divorced category.
There are only approx 27% observations that fall under single category.
There are very few persons with unknown marital status."""
```

```
[156]: 'Majority of the customers i.e. 44% fall into the married category which makes
sense as these may be the persons who require loans to help them supply family
needs.\nThere are only approx 27% customers that lie in divorced
category.\nThere are only approx 27% observations that fall under single
category.\nThere are very few persons with unknown marital status.'
```

```
[158]: ###Bivariate Analysis
# Checking variable distribution in the data
```

```
[160]: """sns.pairplot(data, hue="Loan Status")
plt.show()"""
```

```
[160]: 'sns.pairplot(data, hue="Loan Status")\nplt.show()'
```

```
[162]: - There are overlaps i.e., no clear distinction in the distribution of
↳ variables for people who have defaulted and did not default.
- Let's explore this further with the help of other plots.
```

Cell In[162], line 2

- Let's explore this further with the help of other plots.

SyntaxError: unterminated string literal (detected at line 2)

```
[164]: ### function to plot distributions wrt target
```

```
def distribution_plot_wrt_target(data, predictor, target):
    fig, axs = plt.subplots(2, 2, figsize=(12, 10))

    target_uniq = data[target].unique()

    axs[0, 0].set_title("Distribution of target for target=" +
↳str(target_uniq[0]))
    sns.histplot(
        data=data[data[target] == target_uniq[0]],
        x=predictor,
        kde=True,
        ax=axs[0, 0],
        color="teal",
        stat="density",
    )

    axs[0, 1].set_title("Distribution of target for target=" +
↳str(target_uniq[1]))
    sns.histplot(
        data=data[data[target] == target_uniq[1]],
        x=predictor,
        kde=True,
        ax=axs[0, 1],
        color="orange",
        stat="density",
    )

    axs[1, 0].set_title("Boxplot w.r.t target")
```

```

sns.boxplot(data=data, x=target, y=predictor, ax=axes[1, 0],
palette="gist_rainbow")

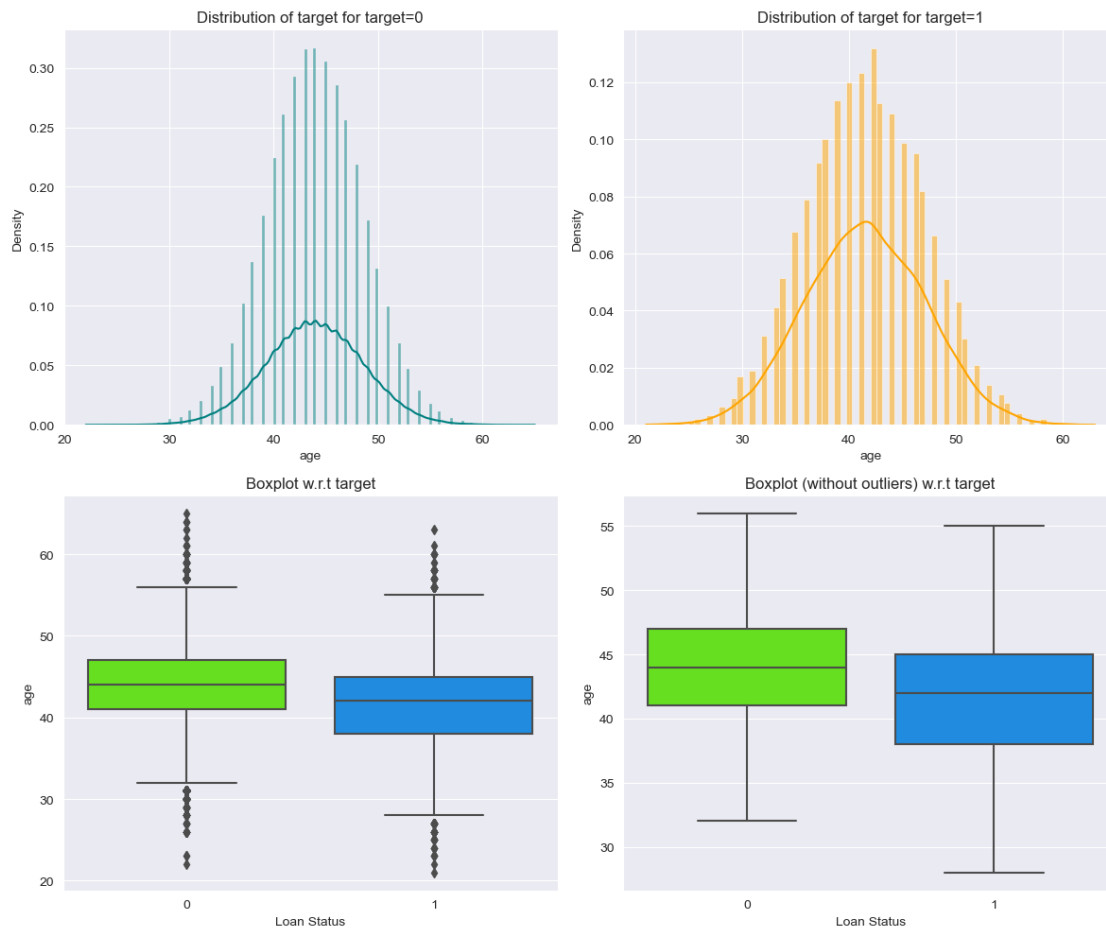
axes[1, 1].set_title("Boxplot (without outliers) w.r.t target")
sns.boxplot(
    data=data,
    x=target,
    y=predictor,
    ax=axes[1, 1],
    showfliers=False,
    palette="gist_rainbow",
)

plt.tight_layout()
plt.show()

```

[166]: *### Loan status vs Age*

[168]: `distribution_plot_wrt_target(data, "age", "Loan Status")`




```
[169]: - We can see that the median age of defaulters is less than the median age of
        ↳ non-defaulters.
        - This shows that younger customers are more likely to default.
        - There are outliers in boxplots of both class distributions
```

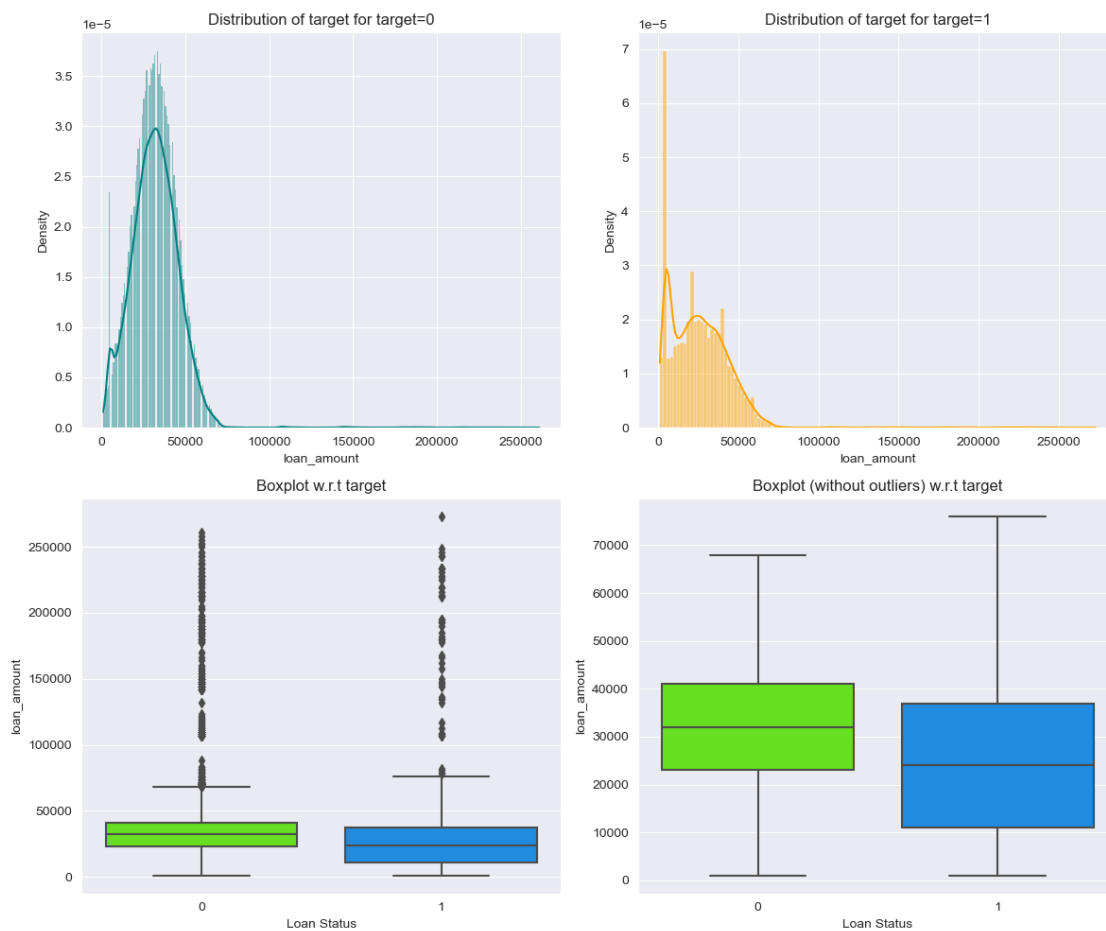
Cell In[169], line 1

```
- We can see that the median age of defaulters is less than the median age
↳ of non-defaulters.
```

SyntaxError: invalid syntax

```
[171]: ### Loan Status vs loan amount
```

```
[173]: distribution_plot_wrt_target(data, "loan_amount", "Loan Status")
```



```
[174]: - We can see that the lower quartile loan amount of defaulters is much less,
      ↪than the lower quartile amount of non-defaulters.
      - This shows that customers with low loan amount are more likely to default.
      - The bank may need to be more cautious when approving smaller loans
      - There are outliers in boxplots of both class distributions
```

```
Cell In[174], line 1
```

```
- We can see that the lower quartile loan amount of defaulters is much less,
↪than the lower quartile amount of non-defaulters.
```

```
SyntaxError: invalid syntax
```

```
[ ]: ### Risk vs remaining term
```

```
[ ]: distribution_plot_wrt_target(data, "remaining term", "Loan Status")
```

```
[178]: -The median remaining term of non defaulters is equal to the median remaining
      ↪term of defaulters
      -This shows that the remaining term may not be a significant factor in
      ↪distinguishing between defaulters and non defaulters
```

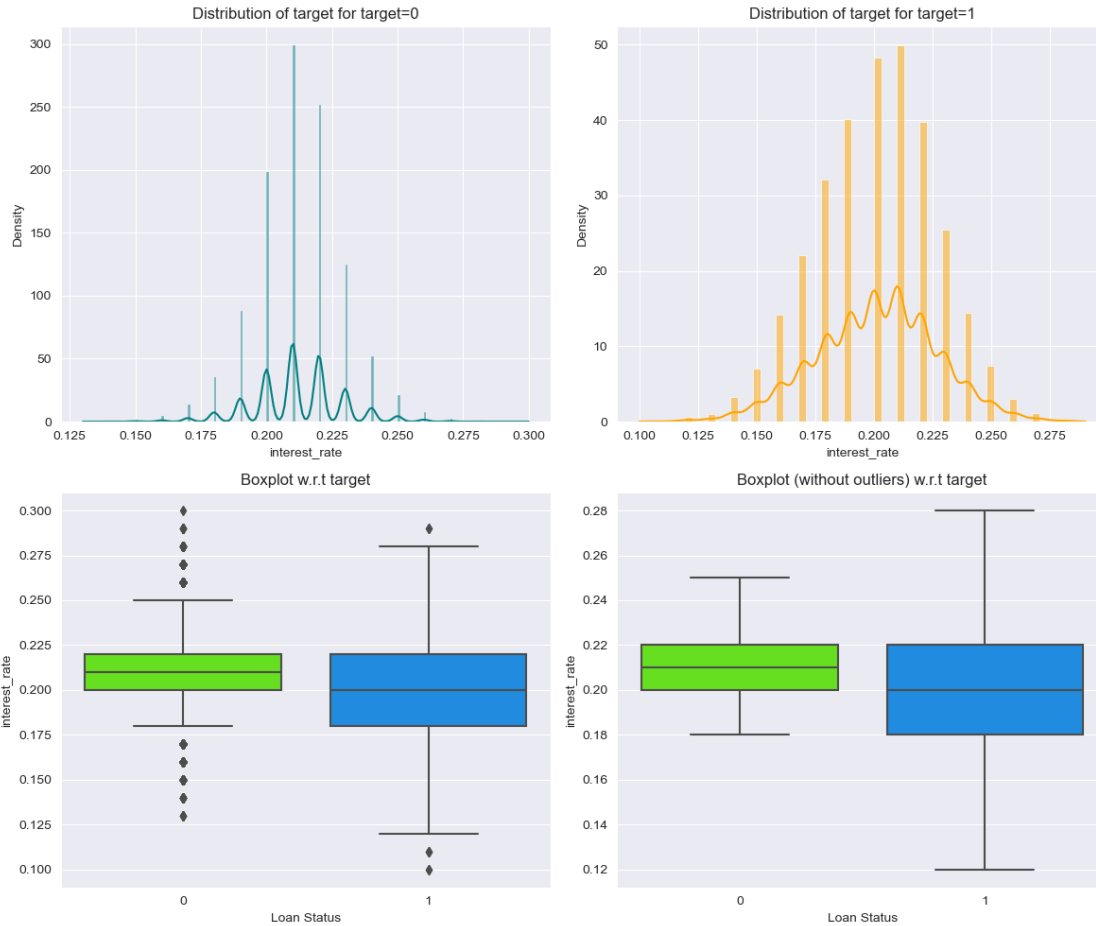
```
Cell In[178], line 1
```

```
-The median remaining term of non defaulters is equal to the median
↪remaining term of defaulters
```

```
SyntaxError: invalid syntax
```

```
[180]: ## Interest vs Loan Status
```

```
[182]: distribution_plot_wrt_target(data, "interest_rate", "Loan Status")
```



[183]: -Lower median interest rate for defaulters compared to non defaulters
 -Lower lower quatile Q1 interest rate for defaulters compared to non defaulters
 -This implies that lower interest rate may not necessarily a guarantee of loan repayment
 ↪repayment
 -The defaulters may be more likely to have lower creditworthiness despite having lower interest rates

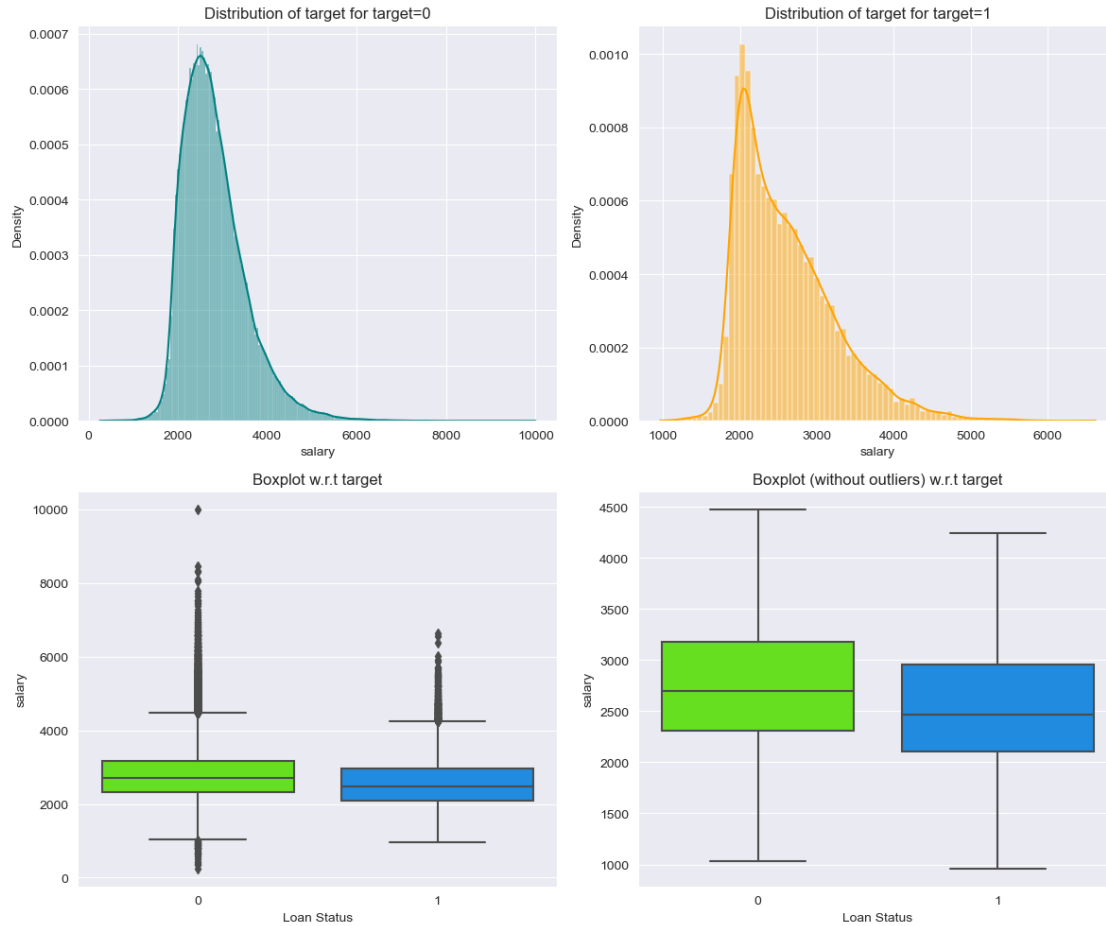
Cell In[183], line 1

-Lower median interest rate for defaulters compared to non defaulters

SyntaxError: invalid syntax

[]: ##Salary vs Loan Status

[186]: distribution_plot_wrt_target(data, "salary", "Loan Status")



[187]: -The median salary for defaulters is much lower than the median salary for non-defaulters
 -This means that lower income customers are more likely to default on loan
 -It also shows that higher income customers tend to have greater financial stability and ability to repay loans

Cell In[187], line 1

-The median salary for defaulters is much lower than the median salary for non defaulters

SyntaxError: invalid syntax

[]: ##Number_of_defaults vs Loan Status

[]: distribution_plot_wrt_target(data, "outstanding_balance", "Loan Status")

```
[ ]: -The median outstanding balance for defaulters is lower than that of non-
      ↳ defaulters
      -This shows that defaulters tend to have lower outstanding balances but still
      ↳ struggle to repay their loans
      -Non defaulters have higher outstanding balance but still manages to repay their
      ↳ loans
```

```
[ ]: # function to plot stacked bar chart
```

```
def stacked_barplot(data, predictor, target):
    """
    Print the category counts and plot a stacked bar chart

    data: dataframe
    predictor: independent variable
    target: target variable
    """
    count = data[predictor].nunique()
    sorter = data[target].value_counts().index[-1]
    tab1 = pd.crosstab(data[predictor], data[target], margins=True).sort_values(
        by=sorter, ascending=False
    )
    print(tab1)
    print("-" * 120)
    tab = pd.crosstab(data[predictor], data[target], normalize="index").
    ↳ sort_values(
        by=sorter, ascending=False
    )
    tab.plot(kind="bar", stacked=True, figsize=(count + 5, 6))
    plt.legend(
        loc="lower left",
        frameon=False,
    )
    plt.legend(loc="upper left", bbox_to_anchor=(1, 1))
    plt.show()
```

```
[193]: ###Loan Status vs Sex
```

```
[195]: stacked_barplot(data, "gender", "Loan Status")
```

```
-----
NameError                                Traceback (most recent call last)
Cell In[195], line 1
----> 1 stacked_barplot(data, "gender", "Loan Status")
```

```
NameError: name 'stacked_barplot' is not defined
```

```
[197]: - We saw earlier that the percentage of male customers is more than the female
      ↪ customers. This plot shows that male customers are more likely to default as
      ↪ compared to female customers.
```

```
Cell In[197], line 1
```

```
- We saw earlier that the percentage of male customers is more than the
↪ female customers. This plot shows that male customers are more likely to
↪ default as compared to female customers.
```

```
SyntaxError: invalid syntax
```

```
[199]: ##Loan amount vs Loan Status
```

```
[201]: stacked_barplot(data, "marital_status", "Loan Status")
```

```
-----
NameError                                Traceback (most recent call last)
```

```
Cell In[201], line 1
```

```
----> 1 stacked_barplot(data, "marital_status", "Loan Status")
```

```
NameError: name 'stacked_barplot' is not defined
```

```
[203]: -This plot shows that divorced customers are more likely to default as compared
      ↪ to single and married customers.
```

```
Cell In[203], line 1
```

```
-This plot shows that divorced customers are more likely to default as
↪ compared to single and married customers.
```

```
SyntaxError: invalid syntax
```

```
[205]: ##Job vs Loan status
```

```
[207]: stacked_barplot(data, "job", "Loan Status")
```

```
-----
NameError                                Traceback (most recent call last)
```

```
Cell In[207], line 1
```

```
----> 1 stacked_barplot(data, "job", "Loan Status")
```

```
NameError: name 'stacked_barplot' is not defined
```

```
[209]: # The plot above shows that lawyers are more likely to default followed by Data
      ↪ Scientists
```

```
[211]: # location vs loan status
      stacked_barplot(data, "location", "Loan Status")
```

```
-----
NameError                                Traceback (most recent call last)
Cell In[211], line 2
      1 # location vs loan status
----> 2 stacked_barplot(data, "location", "Loan Status")

NameError: name 'stacked_barplot' is not defined
```

```
[213]: As we can see from above, customers from above customers from Hwange followed
      ↪ by victoria falls, Gokwe etc are more likely to default
```

```
Cell In[213], line 1
      As we can see from above, customers from above customers from Hwange
      ↪ followed by victoria falls, Gokwe etc are more likely to default
      ~
SyntaxError: invalid syntax
```

```
[215]: # Employment status vs loan status
      stacked_barplot(data, "is_employed", "Loan Status")
```

```
-----
NameError                                Traceback (most recent call last)
Cell In[215], line 2
      1 # Employment status vs loan status
----> 2 stacked_barplot(data, "is_employed", "Loan Status")

NameError: name 'stacked_barplot' is not defined
```

```
[217]: # More customers who are unemployed are likely to default
```

```
[219]: # number of defaults vs loan status
      stacked_barplot(data, "number_of_defaults", "Loan Status")
```

```

NameError                                Traceback (most recent call last)
Cell In[219], line 2
      1 # number of defaults vs loan status
----> 2 stacked_barplot(data, "number_of_defaults", "Loan Status")

NameError: name 'stacked_barplot' is not defined

```

[221]: *# Customers who have 2 records of defaults are more prone to default followed by one*

[223]: *###Model evaluation criterion*
 Model can make wrong predictions as:

Model predicted a non-defaulter as a defaulter - False Positive
 Model predicted a defaulter as a non-defaulter - False Negative

How to reduce this loss i.e need to reduce False Negatives ?

Bank would want to reduce false negatives, this can be done by maximizing the Recall. Greater the recall lesser the chances of false negatives.

```

Cell In[223], line 2
      Model can make wrong predictions as:
      ~
SyntaxError: invalid syntax

```

[225]: First, let's create functions to calculate different metrics and confusion matrix so that we don't have to use the same code repeatedly for each model.

The model_performance_classification_statsmodels function will be used to check the model performance of models.

The confusion_matrix_statsmodels function will be used to plot confusion matrix.

```

Cell In[225], line 1
      First, let's create functions to calculate different metrics and confusion matrix so that we don't have to use the same code repeatedly for each model.
      ~
SyntaxError: invalid syntax

```

[]:


```
[ ]:
```

```
[229]: ### Data Preparation
```

```
[231]: ### Logistic Regression (with statsmodels library)
```

```
[233]: X = data[
    [
        "gender",
        "is_employed",
        "job",
        "remaining term",
        "loan_amount",
        "number_of_defaults",
        "outstanding_balance",
        "interest_rate",
        "age",
        "salary",
        "marital_status",
    ]
]
Y = data["Loan Status"]
```

```
[235]: # creating dummy variables
X = pd.get_dummies(X, drop_first=True)

# standardising continuous variables
scaler = StandardScaler()
X[
    [
        "interest_rate",
        "remaining term",
        "salary",
        "outstanding_balance",
        "age",
        "loan_amount",
    ]
] = scaler.fit_transform(
    X[
        [
            "interest_rate",
            "remaining term",
            "salary",
            "outstanding_balance",
            "age",
            "loan_amount",
        ]
    ]
)
```

```

    ]
)

# Saving the Standard Scaler
pickle.dump(
    scaler,
    open(
        "C:/Users/Munashe Muchinako/OneDrive/Desktop/data science/Data Science_
↪Competition Question and Data/fastapi endpoints/ML Models/StandardScaler.pkl",
        "wb",
    ),
)

# adding constant
X = sm.add_constant(X)

# splitting in training and test set
X_train, X_test, y_train, y_test = train_test_split(X, Y, test_size=0.2,
↪random_state=1)

```

```
[237]: print(X_train.shape, X_test.shape)
```

```
(80000, 22) (20000, 22)
```

```
[239]: # Initialising models
```

```
[241]: """models = {
    "Logistic Regression": LogisticRegression(max_iter=1000),
    "Decision Tree": DecisionTreeClassifier(),
    "Random Forest": RandomForestClassifier(n_estimators=100),
    "Gradient Boosting": GradientBoostingClassifier(n_estimators=100),
    "SVM": SVC(probability=True),
}"""
```

```
[241]: 'models = {\n    "Logistic Regression": LogisticRegression(max_iter=1000),\n    "Decision Tree": DecisionTreeClassifier(),\n    "Random Forest":\nRandomForestClassifier(n_estimators=100),\n    "Gradient Boosting":\nGradientBoostingClassifier(n_estimators=100),\n    "SVM":\nSVC(probability=True),\n}'
```

```
[243]: # Define model parameters
model_params = {
    "Logistic Regression": [
        {"solver": "liblinear", "C": 0.1},
        {"solver": "liblinear", "C": 1.0},
        {"solver": "liblinear", "C": 10.0},
    ],

```

```

        {"solver": "newton-cg", "C": 1.0},
        {"solver": "saga", "C": 1.0},
    ],
    "Decision Tree": [
        {"max_depth": 5, "min_samples_split": 10},
        {"max_depth": 10, "min_samples_split": 5},
        {"max_depth": None, "min_samples_split": 10},
        {"max_depth": 5, "min_samples_split": 2},
        {"max_depth": 15, "min_samples_split": 5},
    ],
    "Random Forest": [
        {"n_estimators": 50, "max_depth": 10},
        {"n_estimators": 100, "max_depth": 15},
        {"n_estimators": 200, "max_depth": None},
        {"n_estimators": 100, "max_depth": 10},
        {"n_estimators": 150, "max_depth": 20},
    ],
    "Gradient Boosting": [
        {"n_estimators": 50, "learning_rate": 0.1},
        {"n_estimators": 100, "learning_rate": 0.1},
        {"n_estimators": 150, "learning_rate": 0.05},
        {"n_estimators": 100, "learning_rate": 0.01},
        {"n_estimators": 200, "learning_rate": 0.1},
    ],
]
}

```

```
[245]: # Train and Evaluate the models
```

```

[247]: for algorithm, params_list in model_params.items():
        print(f"\n{algorithm} Models:")
        for i, params in enumerate(params_list):
            print(f"\nModel {i+1} with parameters: {params}")

            if algorithm == "Logistic Regression":
                model = LogisticRegression(max_iter=1000, **params)
            elif algorithm == "Decision Tree":
                model = DecisionTreeClassifier(**params)
                model.fit(X_train, y_train)
                # Saving the Decision Tree
                pickle.dump(
                    model,
                    open(
                        "C:/Users/Munashe Muchinako/OneDrive/Desktop/data science/
↳Data Science Competition Question and Data/fastapi endpoints/ML Models/
↳DecisionTree.pkl",
                        "wb",
                    ),
                ),

```

```

    )

    elif algorithm == "Random Forest":
        model = RandomForestClassifier(**params)
    elif algorithm == "Gradient Boosting":
        model = GradientBoostingClassifier(**params)
    elif algorithm == "SVM":
        model = SVC(probability=True, **params)

    model.fit(X_train, y_train)
    y_pred = model.predict(X_test)
    y_pred_prob = (
        model.predict_proba(X_test)[:, 1]
        if hasattr(model, "predict_proba")
        else model.decision_function(X_test)
    )

    # Function for Metrics
    accuracy = accuracy_score(y_test, y_pred)
    f1 = f1_score(y_test, y_pred)
    recall = recall_score(y_test, y_pred)
    precision = precision_score(y_test, y_pred)
    roc_auc = roc_auc_score(y_test, y_pred_prob)
    cm = confusion_matrix(y_test, y_pred)

    # Displaying the outcomes
    print(f"{algorithm}:")
    print(f"Accuracy:{accuracy:.4f}")
    print(f"F1 Score:{f1:.4f}")
    print(f"Recall:{recall:.4f}")
    print(f"Precision:{precision:.4f}")
    print(f"ROC AUC:{roc_auc:.4f}")
    """print("Confusion Matrix:")
    #print(cm)"""

    # Plot precision-Recall curve
    precision_vals, recall_vals, _ = precision_recall_curve(y_test, y_pred_prob)
    plt.plot(recall_vals, precision_vals, marker=".", label=f"{algorithm}")
    plt.xlabel("Recall")
    plt.ylabel("Precision")
    plt.title(f"{algorithm} Precision-Recall Curve")
    plt.legend()
    plt.show()

    # Plot Roc Curve
    fpr, tpr, _ = roc_curve(y_test, y_pred_prob)
    plt.plot(fpr, tpr, marker=".", label=f"{algorithm} (AUC={roc_auc:.4f})")

```

```

plt.xlabel("False Positive Rate")
plt.ylabel("True Positive Rate")
plt.title(f"{algorithm} ROC Curve")
plt.legend()
plt.show()

# defining a function to plot the confusion_matrix of a classification model
disp = ConfusionMatrixDisplay(confusion_matrix=cm, display_labels=model.
↪classes_)
    """plt.figure(figsize=(7, 5))
    sns.heatmap(cm, annot=True, fmt="g")
    plt.xlabel("Predicted Values")
    plt.ylabel("Actual Values")"""
disp.plot()
plt.title(f"{algorithm} Confusion Matrix")
plt.show()
    """

    cm = confusion_matrix(y_train, pred_train)
    plt.figure(figsize=(7, 5))
    sns.heatmap(cm, annot=True, fmt="g")
    plt.xlabel("Predicted Values")
    plt.ylabel("Actual Values")
    plt.show()"""

```

Logistic Regression Models:

Model 1 with parameters: {'solver': 'liblinear', 'C': 0.1}

Model 2 with parameters: {'solver': 'liblinear', 'C': 1.0}

Model 3 with parameters: {'solver': 'liblinear', 'C': 10.0}

Model 4 with parameters: {'solver': 'newton-cg', 'C': 1.0}

Model 5 with parameters: {'solver': 'saga', 'C': 1.0}

Logistic Regression:


Accuracy:0.8716


F1 Score:0.3308


Recall:0.2065

Precision:0.8312

ROC AUC:0.7345







Decision Tree Models:

Model 1 with parameters: {'max_depth': 5, 'min_samples_split': 10}

Model 2 with parameters: {'max_depth': 10, 'min_samples_split': 5}

Model 3 with parameters: {'max_depth': None, 'min_samples_split': 10}

Model 4 with parameters: {'max_depth': 5, 'min_samples_split': 2}

Model 5 with parameters: {'max_depth': 15, 'min_samples_split': 5}

Decision Tree:


Accuracy:0.8898


F1 Score:0.5449


Recall:0.4293

Precision:0.7458

ROC AUC:0.7620







Random Forest Models:

Model 1 with parameters: {'n_estimators': 50, 'max_depth': 10}

Model 2 with parameters: {'n_estimators': 100, 'max_depth': 15}

Model 3 with parameters: {'n_estimators': 200, 'max_depth': None}

Model 4 with parameters: {'n_estimators': 100, 'max_depth': 10}

Model 5 with parameters: {'n_estimators': 150, 'max_depth': 20}

Random Forest:

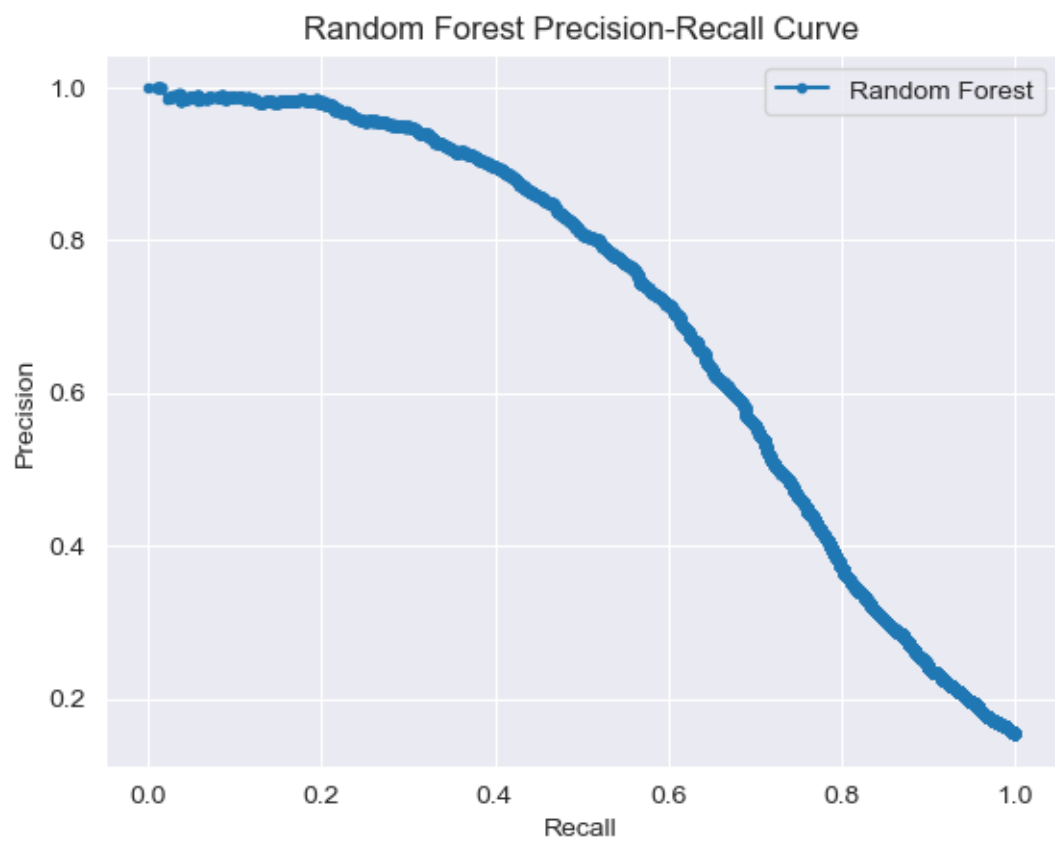
Accuracy:0.9008


F1 Score:0.5552

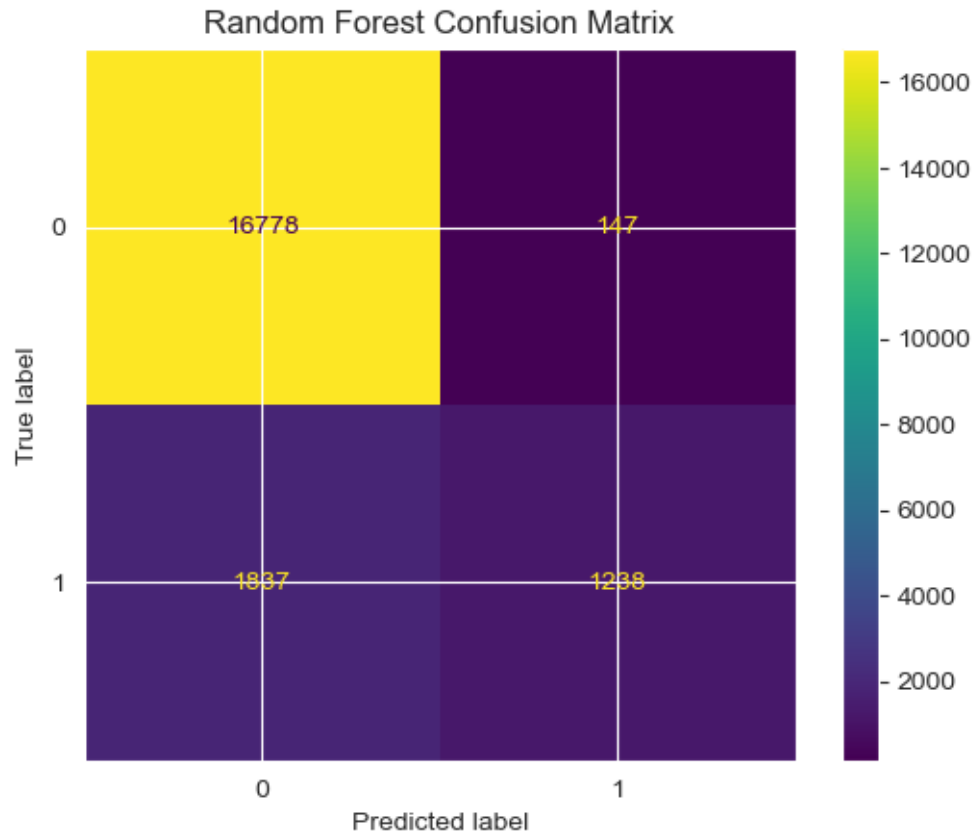
Recall:0.4026

Precision:0.8939

ROC AUC:0.8627







Gradient Boosting Models:

Model 1 with parameters: {'n_estimators': 50, 'learning_rate': 0.1}

Model 2 with parameters: {'n_estimators': 100, 'learning_rate': 0.1}

Model 3 with parameters: {'n_estimators': 150, 'learning_rate': 0.05}

Model 4 with parameters: {'n_estimators': 100, 'learning_rate': 0.01}

Model 5 with parameters: {'n_estimators': 200, 'learning_rate': 0.1}

Gradient Boosting:

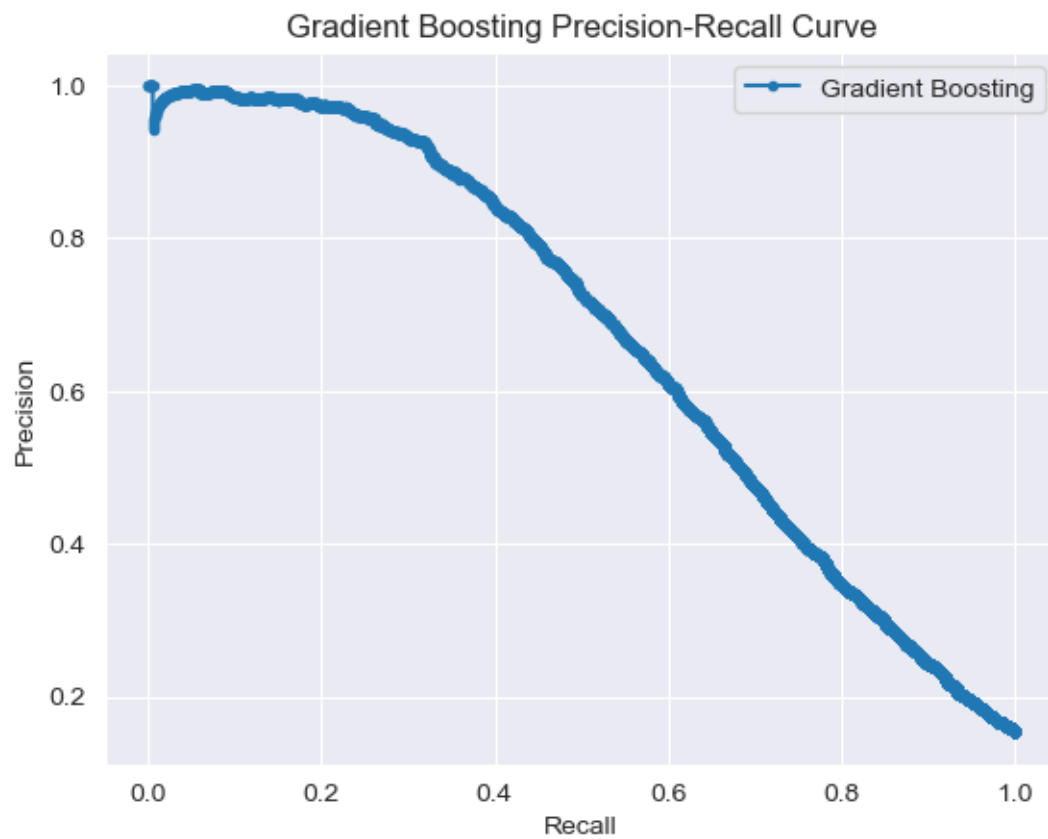
Accuracy:0.8959

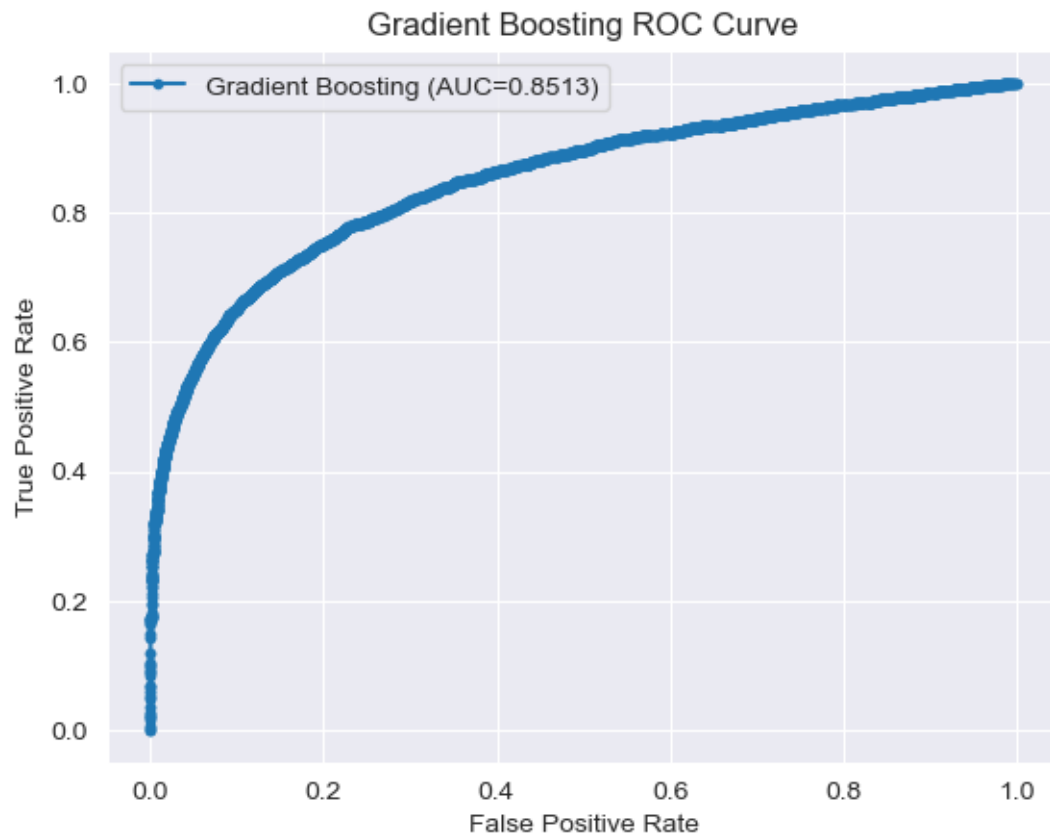
F1 Score:0.5328

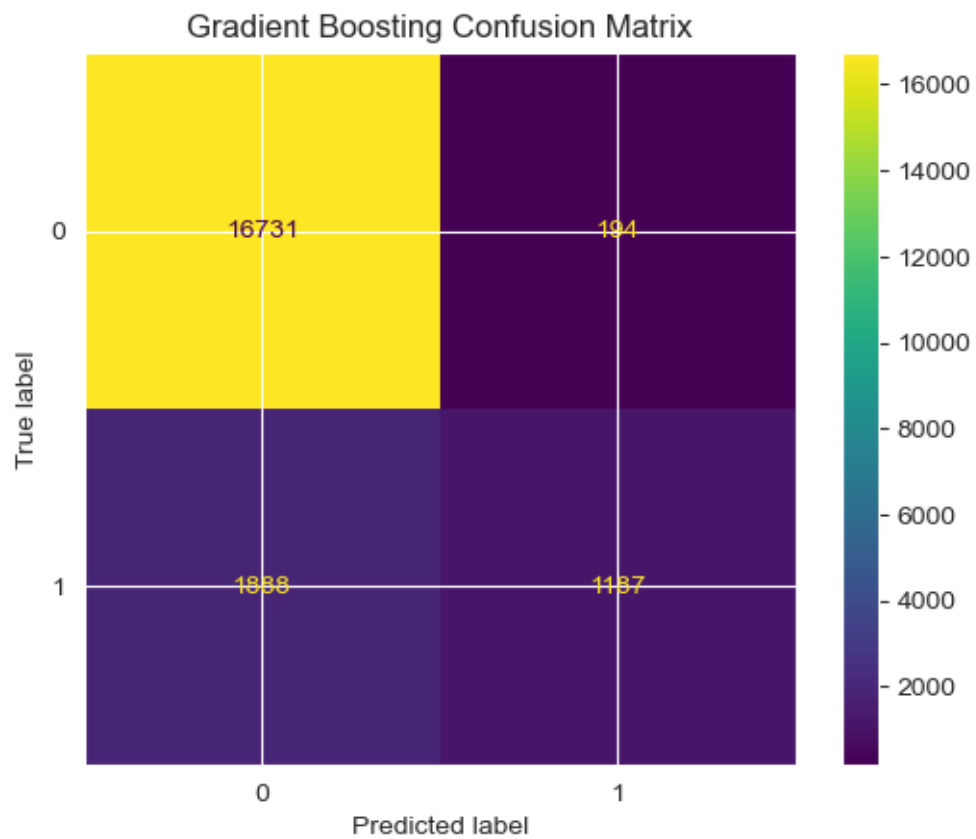
Recall:0.3860

Precision:0.8595

ROC AUC:0.8513







[]:

[]: